

## This is the k-nearest neighbors workbook for ECE C147/C247 Assignment #2

Please follow the notebook linearly to implement k-nearest neighbors.

Please print out the workbook entirely when completed.

The goal of this workbook is to give you experience with the data, training and evaluating a simple classifier, k-fold cross validation, and as a Python refresher.

### ✓ Import the appropriate libraries

```
import numpy as np # for doing most of our calculations
import matplotlib.pyplot as plt # for plotting
from utils.data_utils import load_CIFAR10 # function to load the CIFAR-10 dataset.
```

```
# Load matplotlib images inline
%matplotlib inline
```

```
# These are important for reloading any code you write in external .py files.
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

```
    The autoreload extension is already loaded. To reload it, use:
    %reload_ext autoreload
```

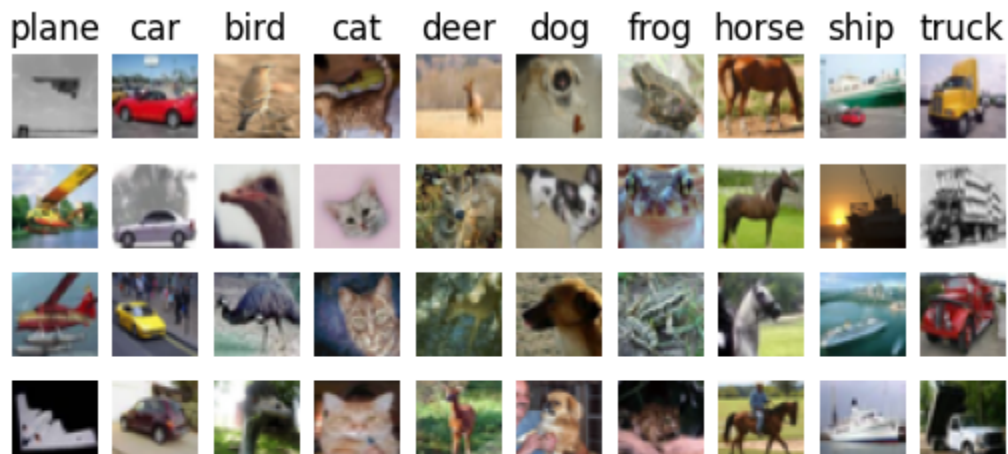
```
# Set the path to the CIFAR-10 data
cifar10_dir = './cifar-10-batches-py' # You need to update this line
X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)
```


```
# As a sanity check, we print out the size of the training and test data.
print('Training data shape: ', X_train.shape)
```

```
print('Training labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

```
Training data shape: (50000, 32, 32, 3)
Training labels shape: (50000,)
Test data shape: (10000, 32, 32, 3)
Test labels shape: (10000,)
```

```
# Visualize some examples from the dataset.
# We show a few examples of training images from each class.
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
num_classes = len(classes)
samples_per_class = 7
for y, cls in enumerate(classes):
    idxs = np.flatnonzero(y_train == y)
    idxs = np.random.choice(idxs, samples_per_class, replace=False)
    for i, idx in enumerate(idxs):
        plt_idx = i * num_classes + y + 1
        plt.subplot(samples_per_class, num_classes, plt_idx)
        plt.imshow(X_train[idx].astype('uint8'))
        plt.axis('off')
        if i == 0:
            plt.title(cls)
plt.show()
```





```
# Subsample the data for more efficient code execution in this exercise
```

```
num_training = 5000
```

```
mask = list(range(num_training))
```

```
X_train = X_train[mask]
```

```
y_train = y_train[mask]
```

```
num_test = 500
```

```
mask = list(range(num_test))
```

```
X_test = X_test[mask]
```

```
y_test = y_test[mask]
```

```
# Reshape the image data into rows
```

```
X_train = np.reshape(X_train, (X_train.shape[0], -1))
```

```
X_test = np.reshape(X_test, (X_test.shape[0], -1))
```

```
print(X_train.shape, X_test.shape)
```

```
(5000, 3072) (500, 3072)
```

## ✓ K-nearest neighbors

In the following cells, you will build a KNN classifier and choose hyperparameters via k-fold cross-validation.

```
# Import the KNN class
```

```
from nn1 import KNN
```

```
from sklearn import knn

# Declare an instance of the knn class.
knn = KNN()

# Train the classifier.
# We have implemented the training of the KNN classifier.
# Look at the train function in the KNN class to see what this does.
knn.train(X=X_train, y=y_train)
```

## Questions

- (1) Describe what is going on in the function `knn.train()`.
- (2) What are the pros and cons of this training step?

## Answers

- (1) `knn.train()` just stores `X_train` and `y_train` as class variables within the KNN class, effectively memorizing the entire training dataset.
- (2) **Pros:** this type of training step is very simple and easy to implement, without any need for explicit computation. Furthermore, it can lead to higher flexibility as updating the parameters of the model is as simple as storing the new dataset. **Cons:** Since the model stores the entirety of the training dataset, it is very memory intensive. Furthermore, since all the computation is done at time of prediction, finding the prediction can be time and compute intensive.

### ✓ KNN prediction

In the following sections, you will implement the functions to calculate the distances of test points to training points, and from this information, predict the class of the KNN.

```
# Implement the function compute_distances() in the KNN class.  
# Do not worry about the input 'norm' for now; use the default definition of the norm  
# in the code, which is the 2-norm.  
# You should only have to fill out the clearly marked sections.
```

```
import time  
time_start =time.time()  
  
dists_L2 = knn.compute_distances(X=X_test)  
  
print('Time to run code: {}'.format(time.time()-time_start))  
print('Frobenius norm of L2 distances: {}'.format(np.linalg.norm(dists_L2, 'fro')))  
  
Time to run code: 15.357426166534424  
Frobenius norm of L2 distances: 7906696.077040902
```

## Really slow code

Note: This probably took a while. This is because we use two for loops. We could increase the speed via vectorization, removing the for loops.

If you implemented this correctly, evaluating `np.linalg.norm(dists_L2, 'fro')` should return: ~7906696

## ✓ KNN vectorization

The above code took far too long to run. If we wanted to optimize hyperparameters, it would be time-expensive. Thus, we will speed up the code by vectorizing it, removing the for loops.

```
# Implement the function compute_L2_distances_vectorized() in the KNN class.  
# In this function, you ought to achieve the same L2 distance but WITHOUT any for loops.  
# Note, this is SPECIFIC for the L2 norm.
```

```
time_start =time.time()  
dists_L2_vectorized = knn.compute_L2_distances_vectorized(X=X_test)
```

```

print('Time to run code: {}'.format(time.time()-time_start))
print('Difference in L2 distances between your KNN implementations (should be 0): {}'.format(np.linalg.norm(dj
    Time to run code: 0.17024898529052734
    Difference in L2 distances between your KNN implementations (should be 0): 0.0

```

## Speedup

Depending on your computer speed, you should see a 10-100x speed up from vectorization. On our computer, the vectorized form took 0.36 seconds while the naive implementation took 38.3 seconds.

## ✓ Implementing the prediction

Now that we have functions to calculate the distances from a test point to given training points, we now implement the function that will predict the test point labels.

```

# Implement the function predict_labels in the KNN class.
# Calculate the training error (num_incorrect / total_samples)
#   from running knn.predict_labels with k=1

error = 1

# ===== #
# YOUR CODE HERE:
#   Calculate the error rate by calling predict_labels on the test
#   data with k = 1. Store the error rate in the variable error.
# ===== #
y_pred = knn.predict_labels(dists_L2_vectorized, k=1)
error = np.mean(y_pred != y_test)
# ===== #
# END YOUR CODE HERE
# ===== #

print(error)

```

0.726

If you implemented this correctly, the error should be: 0.726.

This means that the k-nearest neighbors classifier is right 27.4% of the time, which is not great, considering that chance levels are 10%.

## ✓ Optimizing KNN hyperparameters

In this section, we'll take the KNN classifier that you have constructed and perform cross-validation to choose a best value of  $k$ , as well as a best choice of norm.

### ✓ Create training and validation folds

First, we will create the training and validation folds for use in k-fold cross validation.

```
# Create the dataset folds for cross-validation.
num_folds = 5

X_train_folds = []
y_train_folds = []

# ===== #
# YOUR CODE HERE:
# Split the training data into num_folds (i.e., 5) folds.
# X_train_folds is a list, where X_train_folds[i] contains the
#   data points in fold i.
# y_train_folds is also a list, where y_train_folds[i] contains
#   the corresponding labels for the data in X_train_folds[i]
# ===== #

indices = np.arange(len(X_train))
```

```

indices = np.arange(len(X_train),
np.random.shuffle(indices)

X_train_shuffled = X_train[indices]
y_train_shuffled = y_train[indices]

X_train_folds = np.split(X_train_shuffled, num_folds)
y_train_folds = np.split(y_train_shuffled, num_folds)

# ===== #
# END YOUR CODE HERE
# ===== #

```

## ✓ Optimizing the number of nearest neighbors hyperparameter.

In this section, we select different numbers of nearest neighbors and assess which one has the lowest k-fold cross validation error.

```

time_start =time.time()

ks = [1, 2, 3, 5, 7, 10, 15, 20, 25, 30]

# ===== #
# YOUR CODE HERE:
#   Calculate the cross-validation error for each k in ks, testing
#   the trained model on each of the 5 folds. Average these errors
#   together and make a plot of k vs. cross-validation error. Since
#   we are assuming L2 distance here, please use the vectorized code!
#   Otherwise, you might be waiting a long time.
# ===== #

cross_validation_errors = []

for k in ks:
    fold_errors = []
    for i in range(num_folds):

```



```

for i in range(num_folds):
    X_train_i = np.concatenate([X_train_folds[j] for j in range(num_folds) if j != i])
    y_train_i = np.concatenate([y_train_folds[j] for j in range(num_folds) if j != i])

    X_val_i = X_train_folds[i]
    y_val_i = y_train_folds[i]

    knn.train(X_train_i, y_train_i)
    dists_i = knn.compute_L2_distances_vectorized(X_val_i)
    y_pred_i = knn.predict_labels(dists_i, k=k)

    fold_errors.append(np.mean(y_pred_i != y_val_i))

cross_validation_errors.append(np.mean(fold_errors))

print("Best value for k is :", ks[np.argmin(cross_validation_errors)])
print("Best cross-validation error is:", np.min(cross_validation_errors))


plt.plot(ks, cross_validation_errors, marker='o')
plt.xlabel('k')
plt.ylabel('Cross-validation error')
plt.title('k-NN Cross-validation Error vs. k')
plt.show()


# ===== #
# END YOUR CODE HERE
# ===== #

print('Computation time: %.2f'%(time.time()-time_start))

Best value for k is : 10
Best cross-validation error is: 0.7247999999999999

```





Computation time: 22.19

## Questions:

- (1) What value of  $k$  is best amongst the tested  $k$ 's?
- (2) What is the cross-validation error for this value of  $k$ ?

## ✓ Answers:

- (1) The best value of  $k$  amongst those tested was  $k = 10$ .
- (2) The cross-validation error for  $k = 10$  was  $0.7248 = 72.48\%$

## ✓ Optimizing the norm

Next, we test three different norms (the 1, 2, and infinity norms) and see which distance metric results in the best cross-validation performance.

```
time_start = time.time()

L1_norm = lambda x: np.linalg.norm(x, ord=1)
L2_norm = lambda x: np.linalg.norm(x, ord=2)
Linf_norm = lambda x: np.linalg.norm(x, ord= np.inf)
norms = [L1_norm, L2_norm, Linf_norm]

# ===== #
# YOUR CODE HERE:
# Calculate the cross-validation error for each norm in norms, testing
# the trained model on each of the 5 folds. Average these errors
# together and make a plot of the norm used vs the cross-validation error
# Use the best cross-validation k from the previous part.
#
# Feel free to use the compute_distances function. We're testing just
# three norms, but be advised that this could still take some time.
# You're welcome to write a vectorized form of the L1- and Linf- norms
# to speed this up, but it is not necessary.
# ===== #

k = 10 # best k from previous part
norm_labels = ['L1', 'L2', 'Linf']
cross_validation_errors = []

for norm, label in zip(norms, norm_labels):
    fold_errors = []
    for i in range(num_folds):
        X_train_i = np.concatenate([X_train_folds[j] for j in range(num_folds) if j != i])
        y_train_i = np.concatenate([y_train_folds[j] for j in range(num_folds) if j != i])

        X_test_i = X_train_folds[i]
        y_test_i = y_train_folds[i]

        knn.train(X_train_i, y_train_i)
```

```
.....X_test_i, y_test_i,
dists_i = knn.compute_distances(X_test_i, norm=norm)
y_pred_i = knn.predict_labels(dists_i, k=k)

fold_errors.append(np.mean(y_pred_i != y_test_i))


cross_validation_errors.append(np.mean(fold_errors))


print("Best norm is :", norm_labels[np.argmin(cross_validation_errors)])
print("Best cross-validation error is:", np.min(cross_validation_errors))

plt.plot(norm_labels, cross_validation_errors, marker='o')
plt.xlabel('Norm')
plt.ylabel('Cross-validation error')
plt.title('Cross-validation Error vs. Norm')
plt.show()

# ===== #
# END YOUR CODE HERE
# ===== #
print('Computation time: %.2f'%(time.time()-time_start))

Best norm is : L1
Best cross-validation error is: 0.6942
```





Computation time: 260.12

## Questions:

- (1) What norm has the best cross-validation error?
- (2) What is the cross-validation error for your given norm and  $k$ ?

## Answers:

- (1) The  $L^1$ -norm has best cross-validation error.
- (2) The cross-validation error for  $L^1$ -norm and  $k = 10$  is  $0.6942 = 69.42\%$

## ✓ Evaluating the model on the testing dataset.

Now, given the optimal  $k$  and norm you found in earlier parts, evaluate the testing error of the k-nearest neighbors model.

```
error = 1
```

```
# ===== #
# YOUR CODE HERE:
# Evaluate the testing error of the k-nearest neighbors model
```

```

# Evaluate the testing error of the k-nearest neighbors classifier
# for your optimal hyperparameters found by 5-fold cross-validation.
# ===== #

k = 10 # optimal k value
norm = L1_norm # optimal norm

knn.train(X_train, y_train)

dists = knn.compute_distances(X_test, norm=norm)

y_pred = knn.predict_labels(dists, k=k)

error = np.mean(y_pred != y_test)

# ===== #
# END YOUR CODE HERE
# ===== #

print('Error rate achieved: {}'.format(error))
    Error rate achieved: 0.722

```

## Question:

How much did your error improve by cross-validation over naively choosing  $k = 1$  and using the L2-norm?

## Answer:

Optimal error rate after cross validation was 0.722, down from 0.726 using  $k = 1$  and  $L^2$ -norm. So, the improvement was  $0.004 = 0.4\%$

