

```
import numpy as np
```

```
class Softmax(object):
```

```
    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)
```

```
    def init_weights(self, dims):
        """
        Initializes the weight matrix of the Softmax classifier.
        Note that it has shape (C, D) where C is the number of
        classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims) * 0.0001
```

```
    def loss(self, X, y):
        """
        Calculates the softmax loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """
```

```
    # Initialize the loss to zero.
    loss = 0.0
```

```
    # ===== #
    # YOUR CODE HERE:
    #   Calculate the normalized softmax loss. Store it as the variable loss.
    #   (That is, calculate the sum of the losses of all the training
    #   set margins, and then normalize the loss by the number of
    #   training examples.)
    # ===== #
```

```
    num_train = X.shape[0]
    num_classes = self.W.shape[0]
```

```
    for i in range(num_train):
        scores = X[i].dot(self.W.T)

        sum_exp_scores = np.sum(np.exp(scores))

        correct_class_score = scores[y[i]]
        softmax_probability = np.exp(correct_class_score) / sum_exp_scores

        loss -= np.log(softmax_probability)
```

```
    loss /= num_train
```

```
    # ===== #
    # END YOUR CODE HERE
    # ===== #
```

```
    return loss
```

```
def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the gradient.
```

*Output: grad -- a matrix of the same dimensions as W containing the gradient of the loss with respect to W.*

```
"""

# Initialize the loss and gradient to zero.
loss = 0.0
grad = np.zeros_like(self.W)

# ===== #
# YOUR CODE HERE:
# Calculate the softmax loss and the gradient. Store the gradient
# as the variable grad.
# ===== #

num_train = X.shape[0]
num_classes = self.W.shape[0]

for i in range(num_train):
    scores = X[i].dot(self.W.T)
    exp_scores = np.exp(scores)
    probabilities = exp_scores / np.sum(exp_scores)

    correct_class_prob = probabilities[y[i]]
    loss -= np.log(correct_class_prob)

    for j in range(num_classes):
        grad[j, :] += X[i] * probabilities[j]
    grad[y[i], :] -= X[i]

loss /= num_train
grad /= num_train

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

    for i in np.arange(num_checks):
        ix = tuple([np.random.randint(m) for m in self.W.shape])

        oldval = self.W[ix]
        self.W[ix] = oldval + h # increment by h
        fxph = self.loss(X, y)
        self.W[ix] = oldval - h # decrement by h
        fxmh = self.loss(X, y) # evaluate f(x - h)
        self.W[ix] = oldval # reset

        grad_numerical = (fxph - fxmh) / (2 * h)
        grad_analytic = your_grad[ix]
        rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) +
abs(grad_analytic))
        print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical, grad_analytic,
rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and outputs as loss_and_grad.
    """
```

```

loss = 0.0
grad = np.zeros(self.W.shape) # initialize the gradient as zero

# ===== #
# YOUR CODE HERE:
# Calculate the softmax loss and gradient WITHOUT any for loops.
# ===== #

num_train = X.shape[0]

scores = X.dot(self.W.T)
scores -= np.max(scores, axis=1, keepdims=True)
exp_scores = np.exp(scores)
probabilities = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)

correct_class_probs = probabilities[np.arange(num_train), y]
loss = -np.sum(np.log(correct_class_probs + np.finfo(float).eps)) / num_train

dscores = probabilities
dscores[np.arange(num_train), y] -= 1
dscores /= num_train
grad = np.dot(dscores.T, X)

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
        training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
        means that X[i] has label 0 <= c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training iteration.
    """
    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W

    # Run stochastic gradient descent to optimize W
    loss_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

        # ===== #
        # YOUR CODE HERE:
        # Sample batch_size elements from the training data for use in
        # gradient descent. After sampling,
        # - X_batch should have shape: (batch_size, dim)
        # - y_batch should have shape: (batch_size,)
        # The indices should be randomly generated to reduce correlations
        # in the dataset. Use np.random.choice. It's okay to sample with
    
```

```

# replacement.
# ===== #

indices = np.random.choice(num_train, batch_size)
X_batch = X[indices]
y_batch = y[indices]

# ===== #
# END YOUR CODE HERE
# ===== #

# evaluate loss and gradient
loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
loss_history.append(loss)

# ===== #
# YOUR CODE HERE:
# Update the parameters, self.W, with a gradient step
# ===== #

self.W -= learning_rate * grad

# ===== #
# END YOUR CODE HERE
# ===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
      array of length N, and each element is an integer giving the predicted
      class.
    """
    y_pred = np.zeros(X.shape[1])
    # ===== #
    # YOUR CODE HERE:
    # Predict the labels given the training data.
    # ===== #

    scores = X.dot(self.W.T)

    y_pred = np.argmax(scores, axis=1)

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return y_pred

```